

# Normal distribution and random effects

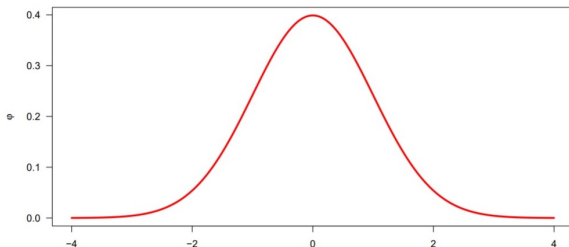
Anders Nielsen and Olav Nikolai Breivik

# The Normal Distribution

- A continuous probability distribution on  $(-\infty, \infty)$
- Probability density function:

$$\phi(x|\mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

- Mean is  $\mu$  and standard deviation is  $\sigma$
- The interval  $(\mu - 2\sigma, \mu + 2\sigma)$  contains 95%



# Complete program using normal likelihood

```

1 library(TMB)
2 compile("norm.cpp")
3 dyn.load(dynlib("norm"))
4
5 data = list()
6 data$Y = rnorm(1000,2,0.3)
7
8 par = list()
9 par$mu = 0
10 par$logSigma = 0
11
12 obj = MakeADFun(data,par,DLL = "norm")
13 opt = nlminb(obj$par,obj$fn,obj$gr)
14 rep = sdreport(obj)

```

```

1 #include <TMB.hpp>
2 template<class Type>
3 Type objective_function<Type>::operator() ()
4 {
5     DATA_VECTOR(Y);
6     PARAMETER(mu);
7     PARAMETER(logSigma);
8
9     Type sigma = exp(logSigma);
10    Type nll = -sum(dnorm(Y,mu,sigma,true));
11    return nll;
12 }

```

- We have now implemented:

$$\text{nll} = - \sum_{i=1}^{1000} \log \phi(y_i | \mu, \sigma),$$

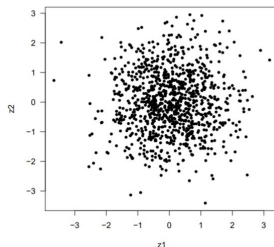
and estimated  $\mu$  and  $\sigma$ .

# Two normal random variables

- Assume we have two independent Gaussian random variables:

$$Z_1 \sim N(0, 1) \text{ and } Z_2 \sim N(0, 1)$$

- Realisations of  $(Z_1, Z_2)$ :



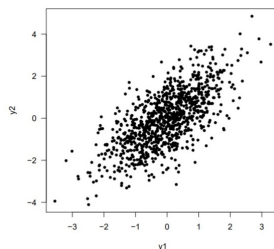
The marginal distribution on each axis is a  $N(0, 1)$ .

# Two normal random variables

- Assume now that

$$\mathbf{Y} = \begin{pmatrix} Z_1 \\ Z_1 + Z_2 \end{pmatrix}$$

- Realisations of  $\mathbf{Y}$ :



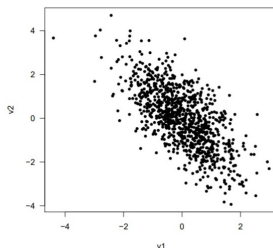
The marginal distributions are  $N(0, 1)$  and  $N(0, 2)$ .

# Two normal random variables

- Assume now that

$$\mathbf{Y} = \begin{pmatrix} Z_1 \\ Z_2 - Z_1 \end{pmatrix}$$

- Realisations of  $\mathbf{Y}$ :



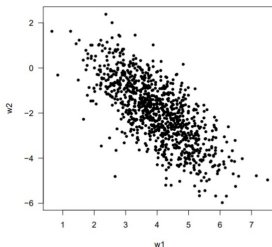
The marginal distributions are  $N(0, 1)$  and  $N(0, 2)$ .

# Two normal random variables

- Assume now that

$$\mathbf{Y} = \begin{pmatrix} Z_1 + 4 \\ Z_2 - Z_1 - 2 \end{pmatrix}$$

- Realisations of  $\mathbf{Y}$ :



The marginal distributions are  $N(4, 1)$  and  $N(-2, 2)$ .

# Two normal random variables

- Note that all these cases can be written as:

$$\mathbf{Y} = \mathbf{AZ} + \mathbf{b},$$

where  $\mathbf{A}$  is a matrix and  $\mathbf{Z} = \begin{pmatrix} Z_1 \\ Z_2 \end{pmatrix}$

- E.g. the last example:

$$\mathbf{Y} = \begin{pmatrix} 1 & 0 \\ -1 & 1 \end{pmatrix} \mathbf{Z} + \begin{pmatrix} 4 \\ -2 \end{pmatrix}.$$



# Multivariate normal distribution

- We say that a  $k$ -dim random variable  $X$  follows a multivariate normal distribution,  $X \sim N(\mu, \Sigma)$ , if there exists a random  $l$ -dim random variable  $Z$  such that  $X \sim AZ + b$ .
- We have that  $\Sigma = AA^t$  and  $\mu = b$ .
- The density of  $X$  is:

$$L(x) = \frac{1}{(2\pi)^{k/2} \sqrt{|\Sigma|}} \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right)$$

- We write  $X \sim N(\mu, \Sigma)$

# Covariance and correlation

- The covariance between two random variables is defined as:

$$\text{cov}(X, Y) = E[(X - \mu_x)(Y - \mu_y)]$$

- When  $X \sim N(\mu, \Sigma)$ , then:

$$\Sigma_{ij} = \text{cov}(X_i, X_j)$$

- We have that the covariance between a variable and itself is the variance:

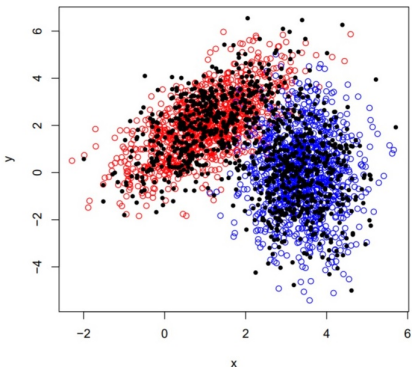
$$\Sigma_{ii} = \text{cov}(X_i, X_i) = \text{var}(X_i)$$

- The correlation coefficient is defined as:

$$\rho_{ij} = \frac{\Sigma_{ij}}{\sqrt{\Sigma_{ii}\Sigma_{jj}}}$$

# Exercise

Assume we have 1000 points from two groups ("red" and "blue"). We are then given 1000 additional points with unknown class. Each of the groups are well described by a 2 dimensional normal distribution. Code to estimate these two normal distributions is provided in `twoNormal.R` and `twoNormal.cpp`. Your task is to assign the most likely class to each of the "black" points.



$$\mathbf{x}_{red} \sim N \left( \begin{pmatrix} \mu_{1,r} \\ \mu_{2,r} \end{pmatrix}, \begin{pmatrix} \sigma_{1,r}^2 & \rho_r \sigma_{1,r} \sigma_{2,r} \\ \rho_r \sigma_{1,r} \sigma_{2,r} & \sigma_{2,r}^2 \end{pmatrix} \right)$$

$$\mathbf{x}_{blue} \sim N \left( \begin{pmatrix} \mu_{1,b} \\ \mu_{2,b} \end{pmatrix}, \begin{pmatrix} \sigma_{1,b}^2 & \rho_b \sigma_{1,b} \sigma_{2,b} \\ \rho_b \sigma_{1,b} \sigma_{2,b} & \sigma_{2,b}^2 \end{pmatrix} \right)$$

# Asymptotically distribution of MLE

- Asymptotically (as we gather more data) the distribution of our estimator will become

$$\hat{\theta} \sim N(\theta_{\text{true}}, H(\theta_{\text{true}})^{-1})$$

where  $H$  is the hessian matrix:

$$H(\theta) = \begin{pmatrix} \frac{\partial^2 f(\theta)}{\partial^2 \theta_1} & \cdots & \frac{\partial^2 f(\theta)}{\partial \theta_1 \partial \theta_k} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f(\theta)}{\partial \theta_k \partial \theta_1} & \cdots & \frac{\partial^2 f(\theta)}{\partial^2 \theta_k} \end{pmatrix}$$

- We can use this result to construct parameter confidence intervals

# Often interested in a function of the parameters

- Say we are interested in an estimate of a differentiable function  $g(\theta)$
- MLE is  $g(\hat{\theta})$
- Uncertainty can be obtained with the delta-method:

$$g(\theta) \stackrel{\text{approx}}{\sim} N(g(\hat{\theta}), \nabla g(\hat{\theta})^T H(\hat{\theta})^{-1} \nabla g(\hat{\theta}))$$

- **Mini Exercise:** Assume we have estimated  $\theta = (\log F_2, \log F_3, \log F_4)$  to  $(-1.13, -0.75, -0.94)$  with an estimated covariance matrix:

$$H(\hat{\theta})^{-1} = \begin{pmatrix} 0.0222 & 0.0135 & 0.0114 \\ 0.0135 & 0.0169 & 0.0137 \\ 0.0114 & 0.0137 & 0.0191 \end{pmatrix}$$

Set up a 95% confidence interval for  $\log \bar{F}_{2-4}$

# Often interested in a function of the parameters

- Say we are interested in an estimate of a differentiable function  $g(\theta)$
- MLE is  $g(\hat{\theta})$
- Uncertainty can be obtained with the delta-method:

$$g(\theta) \stackrel{\text{approx}}{\sim} N(g(\hat{\theta}), \nabla g(\hat{\theta})^T H(\hat{\theta})^{-1} \nabla g(\hat{\theta}))$$

- **Mini Exercise:** Assume we have estimated  $\theta = (\log F_2, \log F_3, \log F_4)$  to  $(-1.13, -0.75, -0.94)$  with an estimated covariance matrix:

$$H(\hat{\theta})^{-1} = \begin{pmatrix} 0.0222 & 0.0135 & 0.0114 \\ 0.0135 & 0.0169 & 0.0137 \\ 0.0114 & 0.0137 & 0.0191 \end{pmatrix}$$

Set up a 95% confidence interval for  $\log \bar{F}_{2-4}$

- **Solution:**  $g(\hat{\theta}) = \log(\frac{1}{3}(e^{\widehat{\log F_2}} + e^{\widehat{\log F_3}} + e^{\widehat{\log F_4}})) \approx -0.928$
- $\nabla g(\hat{\theta}) = \frac{1}{e^{\widehat{\log F_2}} + e^{\widehat{\log F_3}} + e^{\widehat{\log F_4}}} (e^{\widehat{\log F_2}}, e^{\widehat{\log F_3}}, e^{\widehat{\log F_4}})$
- $g(\hat{\theta}) \pm 2 * \sqrt{\nabla g(\hat{\theta})^T H(\hat{\theta})^{-1} \nabla g(\hat{\theta})} \approx -0.928 \pm 2 * 0.122$

Assume  $\mathbf{x} \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ , where

$$\mathbf{x} = \begin{pmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \end{pmatrix}, \boldsymbol{\mu} = \begin{pmatrix} \boldsymbol{\mu}_1 \\ \boldsymbol{\mu}_2 \end{pmatrix}, \boldsymbol{\Sigma} = \begin{pmatrix} \boldsymbol{\Sigma}_{11} & \boldsymbol{\Sigma}_{12} \\ \boldsymbol{\Sigma}_{21} & \boldsymbol{\Sigma}_{22} \end{pmatrix}$$

Then  $\mathbf{x}_1 | \mathbf{x}_2 \sim N(\tilde{\boldsymbol{\mu}}, \tilde{\boldsymbol{\Sigma}})$  where

$$\begin{aligned} \tilde{\boldsymbol{\mu}} &= \boldsymbol{\mu}_1 + \boldsymbol{\Sigma}_{12} \boldsymbol{\Sigma}_{22}^{-1} (\mathbf{x}_2 - \boldsymbol{\mu}_2) \\ \tilde{\boldsymbol{\Sigma}} &= \boldsymbol{\Sigma}_{11} - \boldsymbol{\Sigma}_{12} \boldsymbol{\Sigma}_{22}^{-1} \boldsymbol{\Sigma}_{11}. \end{aligned}$$

Let  $\mathbf{Q} = \boldsymbol{\Sigma}^{-1}$ . We have that

$$\begin{aligned} \text{Var}[x_i | \mathbf{x}_{-i}] &= Q_{ii}^{-1} \\ \text{E}[x_i | \mathbf{x}_{-i}] &= \mu_i - Q_{ii}^{-1} \sum_{i \neq j} Q_{ij} (x_j - \mu_j). \end{aligned}$$

Note that  $Q_{ij} = 0$  is equivalent to that  $x_i$  and  $x_j$  are conditional independent.

# Implement an AR(1) process

- Assume we have  $n = 100$  observations  $x_1, \dots, x_n$  from a mean zero AR(1) process:

$$x_{i+1} = \phi x_i + \epsilon_i, \text{ where } \epsilon_i \sim N(0, \sigma^2)$$

- If the process is in equilibrium when we start observing it, then  $x_1 \sim N(0, \sigma^2/(1 - \phi^2))$
- Exercise a)** Implement the model with use of the univariate normal distribution (`dnorm`)
- Exercise b)** Implement the model with use of multivariate normal function (`AR1_t`)
  - How sparse is the precision matrix to an AR(1) process?
- Verify that you get similar results.



# Latent random effects

- Do not observe directly
- Observe indirectly through the model
  - E.g. fishing mortality in SAM

Simple example: Assume we observe  $Y_1, \dots, Y_n$  where  $Y_i \sim \text{Pois}(\mu_i)$  and

$$\begin{aligned}\mu_i &= \gamma_i \\ \gamma &\sim N(\mathbf{0}, \Sigma).\end{aligned}$$

Here  $\gamma$  is a latent random effect.

# Latent random effects

- Do not observe directly
- Observe indirectly through the model
  - E.g. fishing mortality in SAM

Simple example: Assume we observe  $Y_1, \dots, Y_n$  where  $Y_i \sim \text{Pois}(\mu_i)$  and

$$\begin{aligned}\mu_i &= \gamma_i \\ \gamma &\sim N(\mathbf{0}, \Sigma).\end{aligned}$$

Here  $\gamma$  is a latent random effect.

- If  $\gamma_i$  is Gamma distributed,  $Y_i$  is negative binomial distributed.

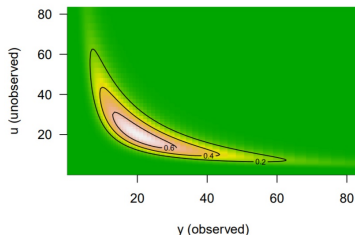
# Latent random effects

- Assume we have:

**Observations:**  $y$

**NOT observed random effects:**  $u$

**Parameters ( $\theta$ ) in the model:**  $(y, u) \sim D(\theta)$



- How do we estimate our parameters when some of our random variables are not observed?

# Latent random effects

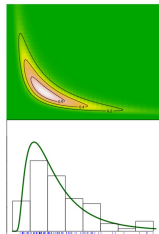
We are typically not interested in the likelihood for a specific  $u$  (it is random)

- 1 Joint model (banana) is determined by model parameters ( $\theta$ )
- 2 Marginal model is calculated by integration
- 3 Marginal is matched to data (what we observe)

The marginal distribution is:

$$L_M(y, \theta) = \int L(y, u, \theta) du.$$

Observations provide information through the marginal distribution



# Latent random effects

- The marginal likelihood is:

$$L_M(\theta, y) = \int_{\mathbb{R}^q} L(\theta, u, y) du$$

- $\theta$  is model parameters
- $y$  is the observed random values (the observations)
- $u$  is the NOT observed random values
- How to calculate the integral?
  - Numerical integration not practical (need a lot of integration points)
  - Seldom an analytical solution, e.g. negative binomial likelihood if...
  - Solution: Approximate with a Taylor-approximation

# Latent random effects

We need to approximate the difficult integral

$$L_M(\theta, y) = \int_{\mathbb{R}^q} L(\theta, u, y) du.$$

Solution:

- Let  $\ell(\theta, u, y) = \log L(\theta, u, y)$ . Note that

$$\ell(\theta, u, y) \approx \ell(\theta, \hat{u}_\theta, y) - \frac{1}{2}(u - \hat{u}_\theta)^t (\ell''_{uu}|_{u=\hat{u}_\theta})(u - \hat{u}_\theta)$$

is a 2.order Taylor approximation around

$$\hat{u}_\theta = \operatorname{argmax}_u \ell(\theta, u, y)$$

for a given  $\theta$ .

Thereby is:

$$\begin{aligned}
 L_M(\theta, y) &= \int_{\mathbb{R}^q} L(\theta, u, y) du \\
 &= \int_{\mathbb{R}^q} \exp(\ell(\theta, u, y)) du \\
 &\approx \int_{\mathbb{R}^q} \exp\left(\ell(\theta, \hat{u}_\theta, y) - \frac{1}{2}(u - \hat{u}_\theta)^t (\ell''_{uu}|_{u=\hat{u}_\theta})(u - \hat{u}_\theta)\right) du \\
 &= L(\theta, \hat{u}_\theta, y) \int_{\mathbb{R}^q} \exp\left(-\frac{1}{2}(u - \hat{u}_\theta)^t (\ell''_{uu}|_{u=\hat{u}_\theta})(u - \hat{u}_\theta)\right) du \\
 &= L(\theta, \hat{u}_\theta, y) \sqrt{\frac{(2\pi)^q}{\det(-\ell''_{uu}|_{u=\hat{u}_\theta})}}
 \end{aligned}$$

The last step is obtained by observing that the integrand has a Gaussian shape.

Taking the logarithm gives:

$$\ell_M(\theta, y) \approx \ell(\theta, \hat{u}_\theta, y) - \frac{1}{2} \log | - \ell''_{uu}|_{u=\hat{u}_\theta}| + \frac{q}{2} \log(2\pi)$$

- This is the Laplace approximation
- Why is automatic differentiation and latent conditional independence structure important to utilize?



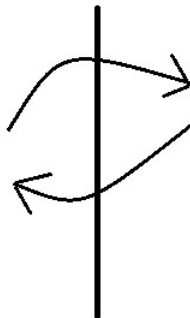
R

C++

Data

Parameters

Optimization routine

 $f(\text{par}, \text{data}) = \dots$

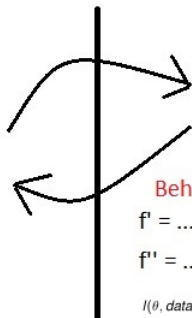
R

C++

Data

Parameters

Optimization routine

 $f(\text{par}, \text{data}) = \dots$ 

Behind the scenes:

 $f' = \dots$  $f'' = \dots$ 

$$l(\theta, \text{data}) \approx f(\theta, \hat{u}_\theta, \text{data}) - \frac{1}{2} \log(|-f''_{uu}|_{u=\hat{u}_\theta}) + \frac{q}{2} \log(2\pi)$$

## Exercise

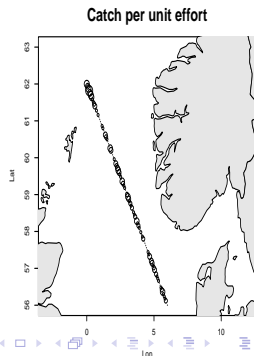
Assume we observe  $Y_1, \dots, Y_n$  where  $Y_i \sim \text{Pois}(\mu_i)$  and

$$\log \mu_i = \gamma_i$$

$$\gamma \sim N(\mathbf{0}, \Sigma),$$

and  $\gamma$  is an AR(1)-process.

- Code is provided in `ar1Latent.R` and `ar1Latent.cpp`
- Data provided in `ar1Latent.RData`
- **Exercise a)** Implement the mode
- **Exercise b)** Inspect the sparseness structure of  $\gamma$ 
  - Why is the sparseness important?



# Solution, R- and C-side

```

1 library(TMB)
2 compile("ar1Latent.cpp")
3 dyn.load("ar1Latent")
4
5 load("cpue.RData")
6 data = list(y = y)
7 par = list(logSigma = -2,
8           phiTrans = 1,
9           gamma = rep(0, length(y)))
10
11 obj = MakeADFun(data, par, random = "gamma",
12               DLL = "ar1Latent")
13 opt = nlminb(obj$par, obj$fn, obj$gr, control
14           = list(trace = 1))
15 rep = sdreport(obj, getJointPrecision =
16               TRUE)

```

```

1 #include <TMB.hpp>
2 template<class Type>
3 Type objective_function<Type>::operator() ()
4 {
5     using namespace density;
6     DATA_VECTOR(y);
7     PARAMETER(logSigma);
8     PARAMETER(phiTrans);
9     PARAMETER_VECTOR(gamma);
10
11     Type phi = Type(2)/(1 + exp(-2*phiTrans))
12             -Type(1);
13     Type sd = exp(logSigma);
14
15     Type nll=0;
16     nll+=SCALE(AR1(phi), sqrt(sd*sd/(1-phi*phi))
17             )(gamma);
18
19     for(int i=1;i<y.size();i++){
20         nll += -dpois(y(i), exp(gamma(i)), true);
21     }
22     return nll;
23 }

```

- Note that  $\gamma$  is now included as random

# Exercise solution

```

outer mgc: 2.125321
  4: 158.72116: -0.405600 1.00234
iter: 1 value: 167.6788 mgc: 0.9840473 ustep: 1
iter: 2 value: 167.6786 mgc: 0.01624102 ustep: 1
iter: 3 value: 167.6786 mgc: 1.878016e-05 ustep: 1
iter: 4 mgc: 3.701173e-11
iter: 1 value: 153.1634 mgc: 0.05411768 ustep: 1
iter: 2 value: 153.1634 mgc: 0.0001761833 ustep: 1
iter: 3 mgc: 1.925322e-09
iter: 1 value: 153.1634 mgc: 0.05411768 ustep: 1
iter: 2 value: 153.1634 mgc: 0.0001761833 ustep: 1
iter: 3 mgc: 1.925322e-09
outer mgc: 0.4393079
  5: 158.64804: -0.348985 1.00601
iter: 1 value: 150.4947 mgc: 0.2225958 ustep: 1
iter: 2 value: 150.4947 mgc: 0.0004934857 ustep: 1
iter: 3 mgc: 9.436551e-09
iter: 1 value: 152.9936 mgc: 0.006990456 ustep: 1
iter: 2 value: 152.9936 mgc: 6.445189e-07 ustep: 1
iter: 3 mgc: 1.110223e-14
iter: 1 value: 152.9936 mgc: 0.006990456 ustep: 1
iter: 2 value: 152.9936 mgc: 6.445189e-07 ustep: 1
iter: 3 mgc: 1.110223e-14
outer mgc: 0.01172901
  6: 158.64298: -0.332983 1.02619
iter: 1 value: 155.2799 mgc: 0.1856167 ustep: 1
iter: 2 value: 155.2799 mgc: 0.0004893153 ustep: 1
iter: 3 mgc: 6.493412e-09
iter: 1 value: 153.249 mgc: 0.01630964 ustep: 1
iter: 2 value: 153.249 mgc: 4.901816e-06 ustep: 1
iter: 3 mgc: 5.8753e-13
iter: 1 value: 153.0547 mgc: 0.001403454 ustep: 1
iter: 2 value: 153.0547 mgc: 5.976854e-08 ustep: 1
iter: 3 mgc: 1.998401e-15
iter: 1 value: 153.0547 mgc: 0.001403454 ustep: 1
iter: 2 value: 153.0547 mgc: 5.976854e-08 ustep: 1
iter: 3 mgc: 1.998401e-15
outer mgc: 0.004273235
  7: 158.64297: -0.332503 1.02580

```

R

Data  
Parameters

Optimization routine

C++

$f(\text{par}, \text{data}) = \dots$

Behind the scenes:

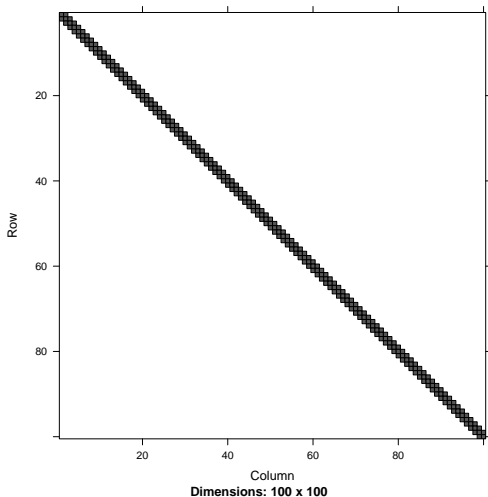
$f^* = \dots$

$f^{\text{st}} = \dots$

$$l(\theta, \text{data}) \approx f(\theta, \hat{\theta}_0, \text{data}) - \frac{1}{2} \log(| - \ell''_{\theta\theta}(\hat{\theta}_0) |) + \frac{Q}{2} \log(2\pi)$$

# Exercise solution

Sparsness structure of AR(1)



# Motivation

Structures in many dimensions can be included in  $\Sigma$

- It is then essential to represent  $\Sigma$  with a sparse precision matrix.

